

1. Downloading the Dataset and importing the Libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn as sk
```

2. Load the dataset

```
In [2]: data = pd.read_csv("Churn_Modelling.csv")
```

```
In [3]: pwd
```

```
Out[3]: 'C:\\Users\\amirt\\Desktop\\IBM Assignments\\Week-2 Assignment\\Team Lead'
```

```
In [4]: #display first five rows
```

```
data.head()
```

```
Out[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bala
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15966
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12551

```
In [5]: #display first 10 rows
```

```
data.head(10)
```

```
Out[5]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bala
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15966
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12551
5	6	15574012	Chu	645	Spain	Male	44	8	11375
6	7	15592531	Bartlett	822	France	Male	50	7	
7	8	15656148	Obinna	376	Germany	Female	29	4	11504
8	9	15792365	He	501	France	Male	44	4	14205
9	10	15592389	H?	684	France	Male	27	2	13460

In [14]: `data.columns`

Out[14]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'], dtype='object')

In [15]: *#unique feature - gender*

`data["Geography"].unique()`

Out[15]: array(['France', 'Spain', 'Germany'], dtype=object)

In [39]: `data.nunique()`

Out[39]:

RowNumber	10000
CustomerId	10000
Surname	2932
CreditScore	460
Geography	3
Gender	2
Age	70
Tenure	11
Balance	6382
NumOfProducts	4
HasCrCard	2
IsActiveMember	2
EstimatedSalary	9999
Exited	2
dtype:	int64

In [45]: `data["IsActiveMember"].unique()`

Out[45]: array([1, 0], dtype=int64)

In [16]: *# To display the bottom of the dataset*

`data.tail()`

Out[16]:


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	
9995	9996	15606229	Obijiaku	771	France	Male	39	5	
9996	9997	15569892	Johnstone	516	France	Male	35	10	5
9997	9998	15584532	Liu	709	France	Female	36	7	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	7
9999	10000	15628319	Walker	792	France	Female	28	4	13

3. Performing (EDA) Visualization

- Univariate Analysis
- Bi - Variate Analysis
- Multi - Variate Analysis


(i) Uni-variate Analysis

```
In [7]: features = ['Age', 'CreditScore', 'Balance']  
data[features].hist(figsize=(13, 10));
```



```
In [8]: sns.displot(data["Age"])
```


```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x27b7b849700>
```



```
In [47]: import warnings  
warnings.filterwarnings("ignore")
```


```
In [48]: sns.boxplot(data["CreditScore"])
```

```
Out[48]: <AxesSubplot:xlabel='CreditScore'>
```




```
In [52]: sns.boxplot(data["NumOfProducts"])
```

```
Out[52]: <AxesSubplot:xlabel='NumOfProducts'>
```




(ii) Bi-variate Analysis


```
In [9]: sns.boxplot(x = data['EstimatedSalary'], y = data['Gender'] );
```



```
In [13]: sns.boxplot(x=data['Gender'],y=data['Balance']);
```



```
In [18]: sns.boxplot(x=data['Geography'],y=data['NumOfProducts']);
```



(iii) Multi-variate Analysis

In [20]: *# Correlation for "NumOfProducts", "EstimatedSalary", "Balance"*


```
df= pd.DataFrame(data,columns=['NumOfProducts','EstimatedSalary','Balance'])
corrMatrix = df.corr()
sns.heatmap(corrMatrix, annot=True)
plt.show()
```



In [21]: *#correlation for all elements in dataset*


```
sns.heatmap(data.corr(),annot = True)
```

Out[21]: <AxesSubplot:>



```
In [28]: sns.relplot(x = "Age",y = "EstimatedSalary",hue="Gender",data=data)
```

```
Out[28]: <seaborn.axisgrid.FacetGrid at 0x27b0df3a520>
```



4. Performing descriptive statistics on the dataset.

```
In [150]: data[['CreditScore','Balance','EstimatedSalary']].mean()
```

```
Out[150]: CreditScore      650.561300
Balance      76485.889288
EstimatedSalary  100090.239881
dtype: float64
```

```
In [151]: #median
```

```
data[['CreditScore','Balance','EstimatedSalary']].median()
```

```
Out[151]: CreditScore      652.000
Balance      97198.540
EstimatedSalary  100193.915
dtype: float64
```

```
In [152]: #mode
```

```
data[['CreditScore','Balance','EstimatedSalary']].mode()
```

```
Out[152]:
```

	CreditScore	Balance	EstimatedSalary
0	850.0	0.0	24924.92

In [153]: *#quantile*

```
data[['CreditScore', 'Balance', 'EstimatedSalary']].quantile()
```

Out[153]: CreditScore 652.000
Balance 97198.540
EstimatedSalary 100193.915
Name: 0.5, dtype: float64

In [154]: *#standard Deviation*

```
data[['CreditScore', 'Balance', 'EstimatedSalary']].std()
```

Out[154]: CreditScore 96.558702
Balance 62397.405202
EstimatedSalary 57510.492818
dtype: float64

In [155]: *#min*

```
data[['CreditScore', 'Balance', 'EstimatedSalary']].min()
```

Out[155]: CreditScore 383.00
Balance 0.00
EstimatedSalary 11.58
dtype: float64

In [156]: *#max*

```
data[['CreditScore', 'Balance', 'EstimatedSalary']].max()
```

Out[156]: CreditScore 850.00
Balance 250898.09
EstimatedSalary 199992.48
dtype: float64

In [157]: *#skew*

```
data[['CreditScore', 'Balance', 'EstimatedSalary']].skew()
```

Out[157]: CreditScore -0.064255
Balance -0.141109
EstimatedSalary 0.002085
dtype: float64

In [26]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                  10000 non-null  int64
8   Balance                 10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary         10000 non-null  float64
13  Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [158]: data.shape

Out[158]: (10000, 14)

In [31]: data.describe()

Out[31]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000

5. Handling the Missing values.

```
In [34]: data.isnull().sum()
```

```
Out[34]: RowNumber      0
CustomerId    0
Surname       0
CreditScore   0
Geography    0
Gender        0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
```


The above result shows that there is no missing values in the dataset

6. Find the outliers and replace the outliers

```
In [55]: def box_scatter(data, x, y):
          fig, (ax1, ax2) = plt.subplots(nrows=2, ncols=1, figsize=(16,6))
          sns.boxplot(data=data, x=x, ax=ax1)
          sns.scatterplot(data=data, x=x, y=y, ax=ax2)
```

```
In [58]: #Scatter and box plot
box_scatter(data, 'CreditScore', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['CreditScore'] < 400])}")
```


of Bivariate Outliers: 19



In the analysis view, there are 19 outliers

```
In [60]: box_scatter(data, 'Age', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['Age'] > 87])}")
```


of Bivariate Outliers: 3



In the above, there are 3 outliers


```
In [61]: box_scatter(data, 'Balance', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['Balance'] > 220000])}")
```

of Bivariate Outliers: 4



Again, there are 4 outliers

```
In [62]: box_scatter(data, 'EstimatedSalary', 'Exited');
plt.tight_layout()
```




Removing the Outliers

```
In [63]: for i in df:
    if df[i].dtype=='int64' or df[i].dtypes=='float64':
        q1=df[i].quantile(0.25)
        q3=df[i].quantile(0.75)
        iqr=q3-q1
        upper=q3+1.5*iqr
        lower=q1-1.5*iqr
        df[i]=np.where(df[i] >upper, upper, df[i])
        df[i]=np.where(df[i] <lower, lower, df[i])
```

```
In [65]: box_scatter(data, 'CreditScore', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['CreditScore'] < 400])}")
```


of Bivariate Outliers: 19



```
In [68]: for i in data:
        if data[i].dtype=='int64' or data[i].dtypes=='float64':
            q1=data[i].quantile(0.25)
            q3=data[i].quantile(0.75)
            iqr=q3-q1
            upper=q3+1.5*iqr
            lower=q1-1.5*iqr
            data[i]=np.where(data[i] >upper, upper, data[i])
            data[i]=np.where(data[i] <lower, lower, data[i])
```


```
In [69]: box_scatter(data, 'CreditScore', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['CreditScore'] < 400])}")
```

of Bivariate Outliers: 19



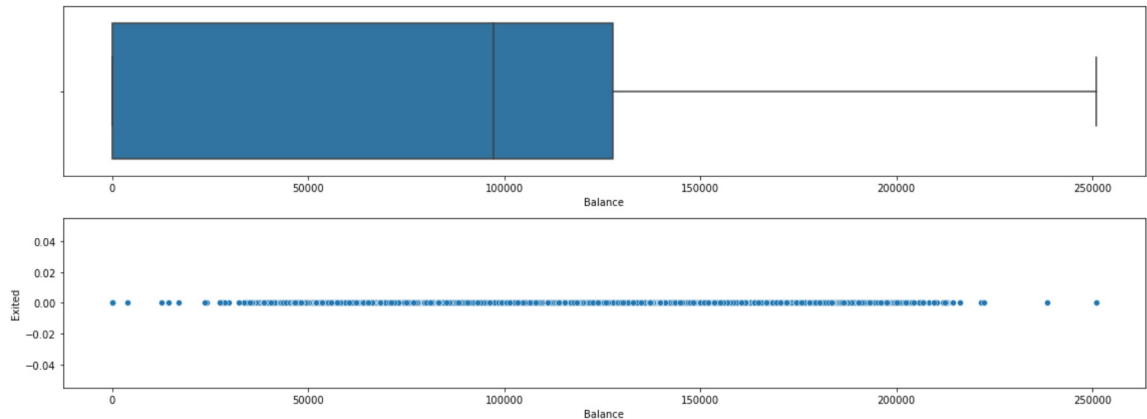
```
In [70]: box_scatter(data, 'Age', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['Age'] > 87])}")
```

of Bivariate Outliers: 0



```
In [71]: box_scatter(data, 'Balance', 'Exited');
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(data.loc[data['Balance'] > 220000])}")
```

of Bivariate Outliers: 4



7. Checking for Categorical columns and performing encoding.

```
In [77]: from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
for i in data:
    if data[i].dtype=='object' or data[i].dtype=='category':
        data[i]=encoder.fit_transform(data[i])
```

8. Split the data into dependent and independent variables

```
In [119]: x = data.iloc[:, :-1]
x.head()
```

Out[119]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bala
0	1.0	15634602.0	1115	619.0	0	0	42.0	2.0	
1	2.0	15647311.0	1177	608.0	2	0	41.0	1.0	8380
2	3.0	15619304.0	2040	502.0	0	0	42.0	8.0	15966
3	4.0	15701354.0	289	699.0	0	0	39.0	1.0	
4	5.0	15737888.0	1822	850.0	2	0	43.0	2.0	12551

```
In [120]: y=data.iloc[: -1]
y.head()
```

```
Out[120]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1.0	15634602.0	1115	619.0	0	0	42.0	2.0	
1	2.0	15647311.0	1177	608.0	2	0	41.0	1.0	8380
2	3.0	15619304.0	2040	502.0	0	0	42.0	8.0	15966
3	4.0	15701354.0	289	699.0	0	0	39.0	1.0	
4	5.0	15737888.0	1822	850.0	2	0	43.0	2.0	12551

9. Scaling the independent variables

```
In [144]: #scaling
names=X.columns
names
```

```
Out[144]: Index(['RowNumber', 'CustomerId', 'Surname', 'Geography', 'Gender', 'Age',
                'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
                'EstimatedSalary', 'Exited'],
                dtype='object')
```

```
In [145]: from sklearn.preprocessing import scale
x= scale(X)
x
```

```
Out[145]: array([[ -1.73187761, -0.78321342, -0.46418322, ...,  0.97024255,
                   0.02188649,  0.          ],
                 [ -1.7315312 , -0.60653412, -0.3909112 , ...,  0.97024255,
                   0.21653375,  0.          ],
                 [ -1.73118479, -0.99588476,  0.62898807, ..., -1.03067011,
                   0.2406869 ,  0.          ],
                 ...,
                 [  1.73118479, -1.47928179,  0.07353887, ...,  0.97024255,
                   -1.00864308,  0.          ],
                 [  1.7315312 , -0.11935577,  0.98943914, ..., -1.03067011,
                   -0.12523071,  0.          ],
                 [  1.73187761, -0.87055909,  1.4692527 , ..., -1.03067011,
                   -1.07636976,  0.          ]])
```

```
In [146]: X = pd.DataFrame(x,columns = names)
X
```

```
Out[146]:
```

	RowNumber	CustomerId	Surname	Geography	Gender	Age	Tenure	Bala
0	-1.731878	-0.783213	-0.464183	-0.901886	-1.095988	0.342615	-1.041760	-1.225
1	-1.731531	-0.606534	-0.390911	1.515067	-1.095988	0.240011	-1.387538	0.117
2	-1.731185	-0.995885	0.628988	-0.901886	-1.095988	0.342615	1.032908	1.333
3	-1.730838	0.144767	-1.440356	-0.901886	-1.095988	0.034803	-1.387538	-1.225
4	-1.730492	0.652659	0.371354	1.515067	-1.095988	0.445219	-1.041760	0.785
...
9995	1.730492	-1.177652	0.580534	-0.901886	0.912419	0.034803	-0.004426	-1.225
9996	1.730838	-1.682806	-0.203004	-0.901886	0.912419	-0.375612	1.724464	-0.306
9997	1.731185	-1.479282	0.073539	-0.901886	-1.095988	-0.273008	0.687130	-1.225
9998	1.731531	-0.119356	0.989439	0.306591	0.912419	0.342615	-0.695982	-0.022
9999	1.731878	-0.870559	1.469253	-0.901886	-1.095988	-1.093840	-0.350204	0.859

10000 rows × 13 columns

10. Splitting the data into Training and Testing

```
In [125]: from sklearn.model_selection import train_test_split
```

```
In [147]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_s
```

```
In [148]: X_train.head()
```

```
Out[148]:
```

	RowNumber	CustomerId	Surname	Geography	Gender	Age	Tenure	Bala
7389	0.827747	-0.195066	0.366627	1.515067	-1.095988	-0.478216	-0.004426	-1.225
9275	1.481077	0.810821	-1.292630	0.306591	0.912419	0.342615	-1.387538	-0.012
2995	-0.694379	-1.507642	0.391445	-0.901886	-1.095988	-0.991236	-1.041760	0.575
5316	0.109639	1.243462	-0.744271	1.515067	0.912419	0.137407	-0.004426	0.467
356	-1.608556	-1.100775	1.117074	1.515067	-1.095988	1.881674	1.032908	0.806

```
In [149]: X_train.shape,y_train.shape,X_test.shape,y_test.shape
```

```
Out[149]: ((8000, 13), (8000,), (2000, 13), (2000,))
```